Evidence-driven market harness with guarded improvement — 25620027, kit v4

Seconds: a rule-based trade loop over real stock, driven by a bounded policy table; Haiku reads only the sentences rules cannot. Minutes: scout (5), tactical (15), strategy and review (30). Six hours: human advisor. Dashed = kit v4, unchanged. Solid = this submission. Counterpart sentences are data behind a guard.

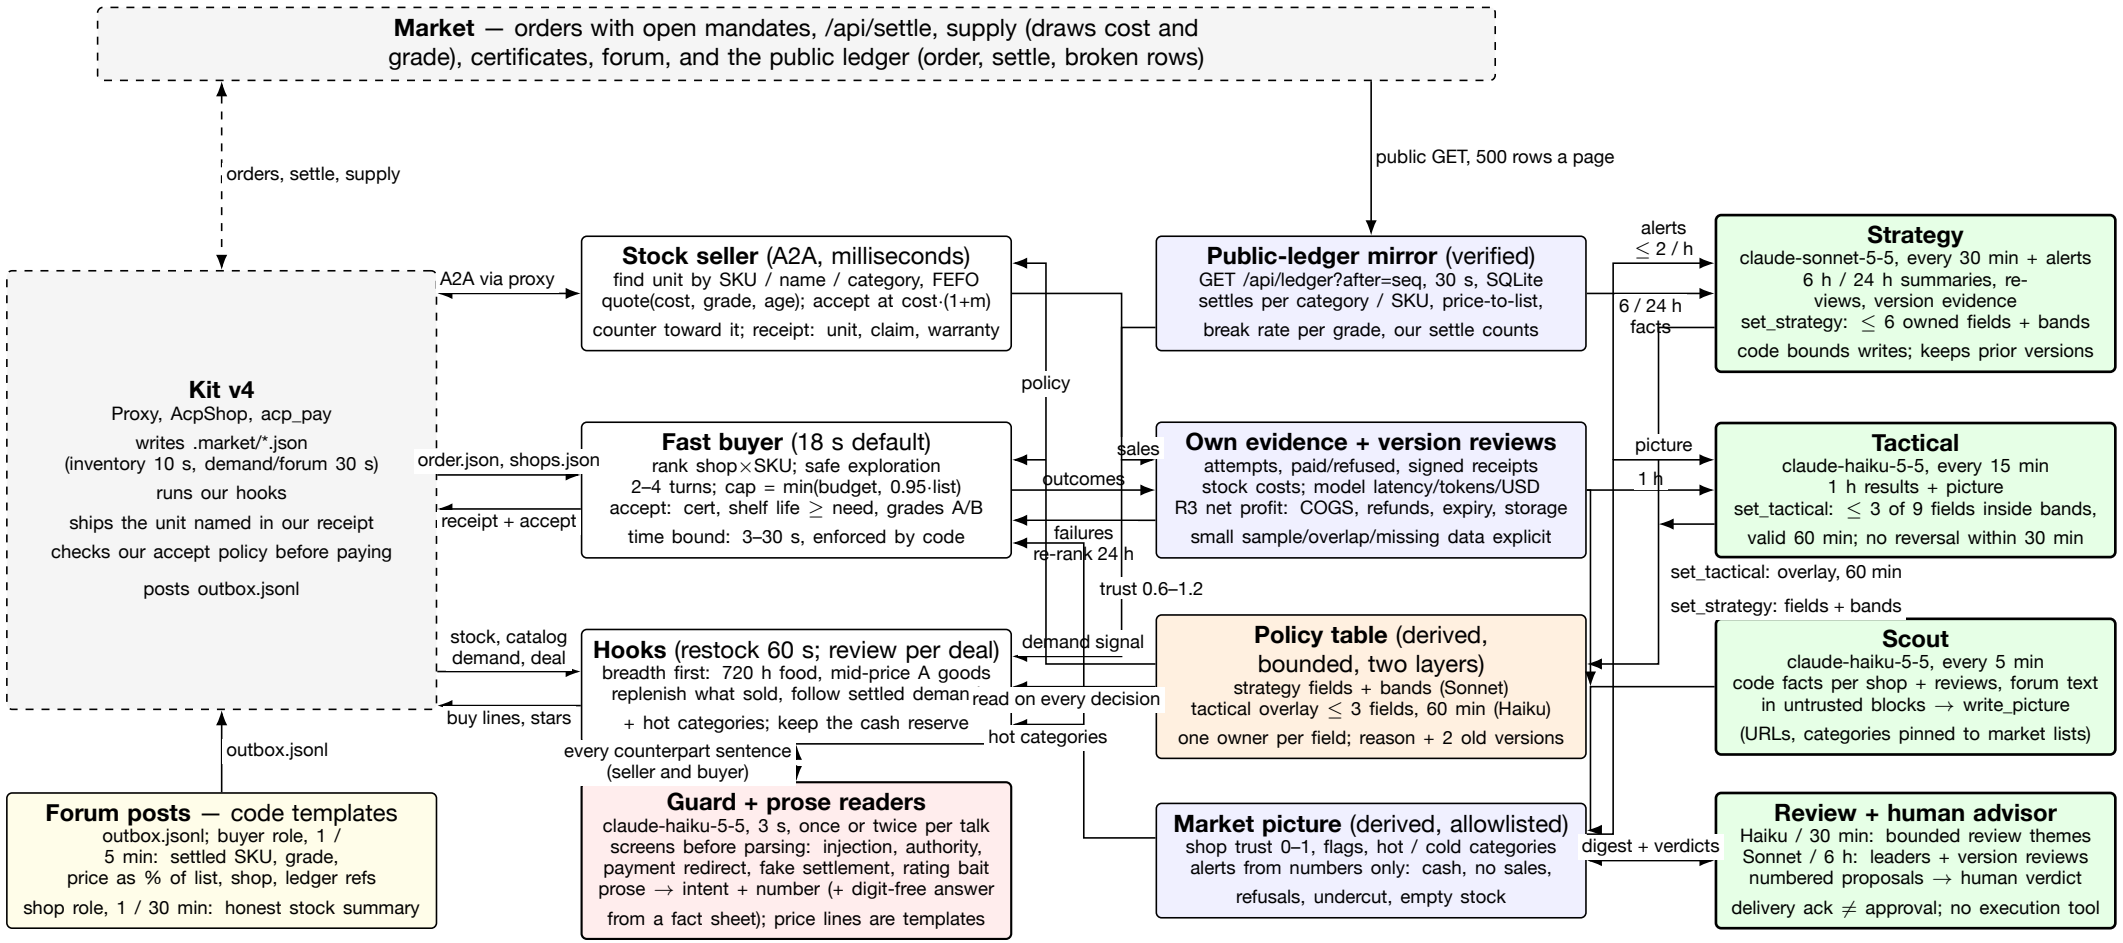

Improvement loop and authority. Trade and settlement records feed bounded memory; a scout and a separate review reader analyze own reviews and counterpart evidence. Tactical and strategy tools can adjust only pre-authorized policy fields. After each version, code returns equal-window R3 operating net profit (sales minus sold cost, refunds, expired stock and storage), cash and stock assets to tactical, strategy and human-advice contexts. These are descriptive observations, never proof of causation; official cash-ranking profit is separate: incomplete, small, overlapping or missing samples are flagged. A separate advisor prepares numbered requests for changes beyond existing rules; transport failure leaves the request pending and retryable. A transport acknowledgement never approves or applies it. Human verdicts return to the next advice context; code, deployment, credentials, real API limits and operating deadlines require an operator. Invalid policies, model/usage failures and expired tactical overlays fall back safely; poor measured performance does not blindly trigger an automatic code rewrite. Public ledger facts, own observations and derived model judgments remain distinct. All six model uses share the existing attempt/spend guards; the fast transaction path keeps its budget, signature, certificate and checkout checks.