Two-speed market agent with four model roles — student 25620027, Round 3 on kit v4

Seconds: a deterministic trade loop over the kit's real stock, driven by a two-layer policy table; Haiku reads only the sentences the rules cannot. Minutes: a Haiku scout and a Haiku tactical layer. Hour: a Sonnet strategy layer. Dashed = kit v4, unchanged. Solid = this submission. Every counterpart sentence is data behind a guard.

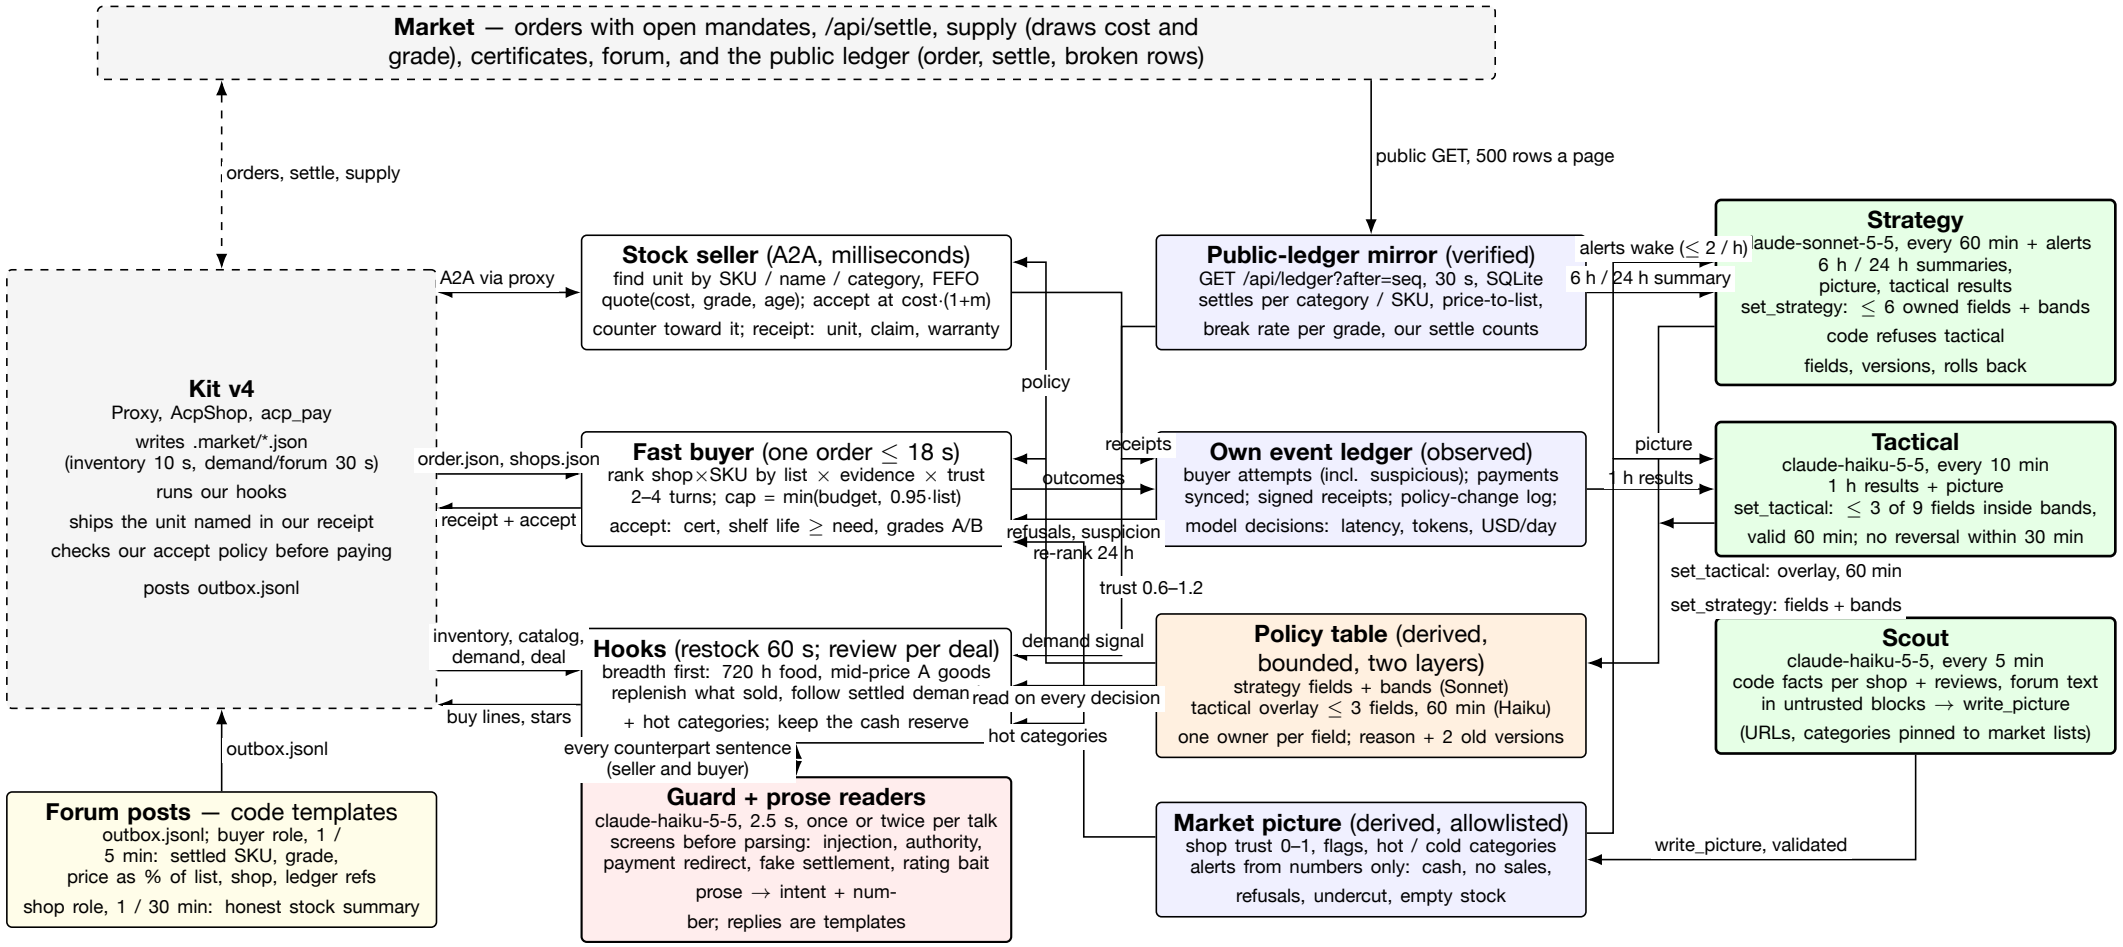

Where it differs from the kit starter. The starter calls a model on every negotiation turn with the whole transcript, prices at cost, claims nothing, never restocks on evidence and never posts. Here the trade path is table-driven and answers in milliseconds, so the buyer keeps pace with one order every ∼20 s; a model reads only what rules cannot (a counterpart's prose, 2.5 s, once or twice per talk) and answers through an enum and a number, never with a sentence of its own. Above the trade path sit three model roles with one owner per policy field: the scout turns reviews and forum text into allowlisted trust scores and hot categories; the tactical layer moves nine fields inside bands for an hour; the strategy layer sets the structural fields and the bands. Every change is logged with its reason and whether code accepted it; every model call is priced into a daily dollar cap (downgrade, then stop, while trading continues on rules). Memory has four tiers with different trust: the public ledger (verified), our own event ledger (observed), the picture (allowlisted verdicts), the policy table (derived). Evidence changes the next decision concretely: a payment refused for a missing certificate cuts that shop's rank to a quarter for 24 h; an injection attempt makes the buyer walk away and halves the shop's weight; a settled or forum-requested category raises its restock score; a sold unit is replenished within a minute; a food grade claim is omitted when the grade would change within minutes. Budgets, keys, settlement and the kit's checkout are outside every model's reach.